

1) FCFS

2) Shortest Job First (Preemptible / Non-preemptible)

Starvation of CPU-bound jobs

Shortest Job Remaining First


3) Round-Robin — Performance depends a lot on predecided time quanta.

### Types of Processes

1) I/O Bound or interactive → Desktops or mobiles

2) CPU Bound or Compute-intensive

Server systems




Large → Close to FCFS  
Small → Repeated process switching

Playing music,  
bash shell,  
file manager

4) Multilevel Queue : —

$P_1$   
0.6


$Q_1$



→ Round robin with some time quanta  $q_1$

0.3  $p_2$


$Q_2$



$q_2$

0.1  $p_3$

$Q_3$



$p_1 > p_2 > p_3$

$q_3 \checkmark$   
(FCFS)  $\infty$

$q_1 < q_2 < q_3$

Interactive processes or processes with high priority go to higher queues.

One possible way:  $\rightarrow$  If there is some process in

both  $Q_i$  and  $Q_j$ , ( $i < j$ ), then execute always with high prob. from  $Q_i$

x 1) Somebody has to predict where to place the process.


2) Starvation of CPU-intensive processes is still possible

$\Downarrow$   
Method 2 problem eliminated

0.4 random (0, 1)

5) Multilevel Feedback Queue

3)  $\checkmark$  Somebody needs to decide which queue it should be placed in.



All processes will start in Q<sub>1</sub>. If it does not finish in one time quantum, it will be pushed down to Q<sub>2</sub> for scheduling again.

## Linux 2.4

Priorities assigned to processes with values ranging from 0 to 139

Priorities from 0 to 99 was for real-time processes,

if some type of performance guarantee is required  
 process has to be executed within a deadline

Soft real-time

You can occasionally miss the deadline, but the more you miss, the


Hard real-time

You cannot miss the deadline

worse is the performance  
of the scheduler,

Cannot be done on  
desktop/mobile system  
↓  
By specialized hardware

Normal processes  $\rightarrow$  priority values from 100 to 139.



Static priority  $\rightarrow$  whatever priority a process starts with  
100 to 139

↓  
 $120 + n$ , where  $n$  is the "nice number"  
 $n = 19$  to  $-20$

$n$  is greater, that process has lower priority

By default,  $n$  is 0.

Root users can make their nice numbers negative.  
Non-root users cannot.

All users can set positive nice numbers for their processes.

Dynamic priority:  $\rightarrow \max(100, \text{MIN}(\text{static priority} - \text{bonus} + 5, 139))$

Distinishing between I/O  
and CPU oriented processes.

CPU-bound process should get less priority;  
so if your waiting time is higher, it should  
get higher bonus.

$$\text{bonus} = \left\lfloor \frac{\text{Waiting time in ms}}{100} \right\rfloor \quad \text{if waiting time} < 1s$$

$$= 10 \quad \text{if} \quad \text{priority} < 120; \quad \text{Time slice} = (140 - \text{prio}) \times 20 \text{ ms}$$
$$\geq 120; \quad \text{Time slice} = (140 - \text{prio}) \times 5 \text{ ms}$$