

O(1) Scheduler — 2.4 Big-oh

↳ Multi-level feedback queue scheduler

→ low time-complexity, so suitable for a lot of processes

→ could handle interactive and CPU-intensive processes.

Disadvantages: —

- (1) A lot of heuristics are involved.
- (2) Abrupt changes in the time slices.

Completely Fair Scheduler


Target assigning of equal processing time to all the processes.

Suppose there are n processes.

$$t_s \propto \frac{1}{n}$$

$t_s = \frac{k}{n}$, k is a constant of proportionality, and here it is known as "target latency".

Process runtime r_i of process p_i



How much Time
in running state

$$\frac{20}{n=1}$$

$$\frac{20}{1}$$

$\left(\begin{array}{l} 2 \text{ ms in running state} \\ 30 \text{ ms in waiting state} \end{array} \right) \rightarrow \text{least amount of runtime.}$
 \searrow Scheduled for execution

$$n=5$$

$$\frac{20}{5} = 4 \text{ ms} \checkmark$$

If n becomes large, then the time slice will gradually become small.

No process will receive a time slice smaller than 1 ms.

$$\text{CFS} \rightarrow \max \left(1, \frac{k}{n} \right) \text{ in milliseconds}$$

is not in practice completely fair

How is it implemented?

→ Keep track of the run time

→ Using vruntime variable in sched-entity structure within the Linux kernel

Process with minimum vruntime

— Using a height-balanced tree

→ AVL Tree } → $O(\log n)$
→ Red-black Tree ✓

→ Rebalancing happens ~~too~~ much more often

What are the system calls related to scheduling?

nice numbers → CFS also takes care of nice numbers.

How?

$$\underline{\text{vruntime}} = \text{runtime} * \left[\text{nice_number} + \text{21} \right]$$

① nice(1) ✓

Increments the nice value by the given number.

vruntime → unsigned long

② sched_yield();

Soft real-time jobs / processes

→ SCHED_FIFO, SCHED_RR
real-time processes

